

# SQL HINTS, TIPS, TRICKS AND TUNING

Susanne Ebrecht

PGCon 2013, Ottawa  
May 2013



# SPEAKER

- Susanne Ebrecht
- Diploma in Computer Sciences
- Open Source activity 1996
- Expert in Databases and  
Localisation / Globalisation
- Member of SQL Standard Committee







# RULES

- Interposed Questions are Welcome
- Twitter @miraceesusanne
- No Individual Consulting
- Slides have no informative value outside the talk

# NICKS





# WORK\_MEM

- `SHOW work_mem;`
- `SET work_mem='64MB';`
- Per Session
- $(\text{total RAM} / \text{max\_connections}) / \text{query execution steps}$

# SQL

Data Definition Language

CREATE, ALTER, DROP

Data Modification Language

INSERT, UPDATE, DELETE

Data Query Language

SELECT

Data Control Language

GRANT, REVOKE

Transaction Control Language

START TRANSACTION,  
SAVEPOINT, COMMIT,  
ROLLBACK

# SQL TUNING

- DML Tuning thwarts DQL
- DQL Tuning thwarts DMS

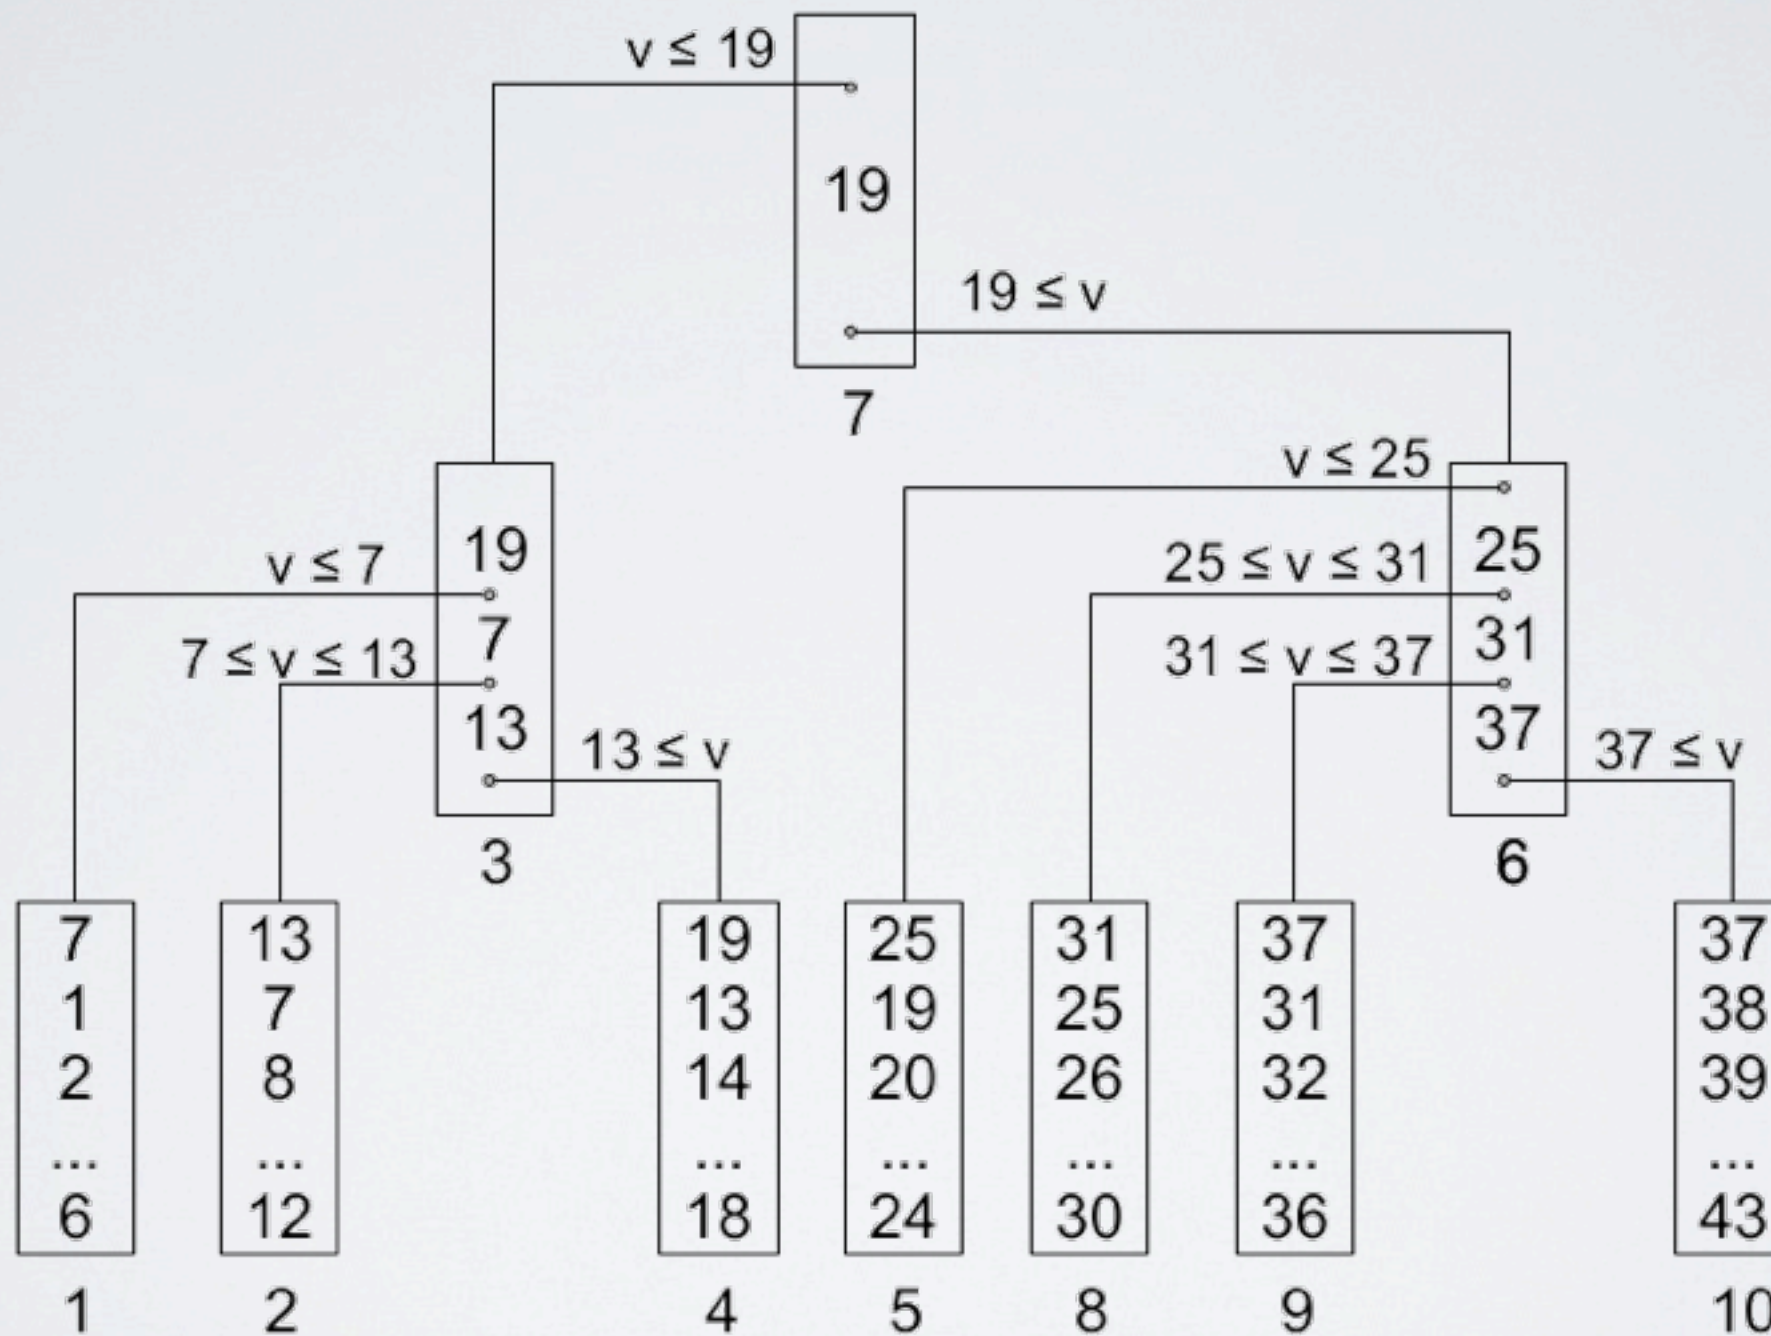


# NORMALIZATION

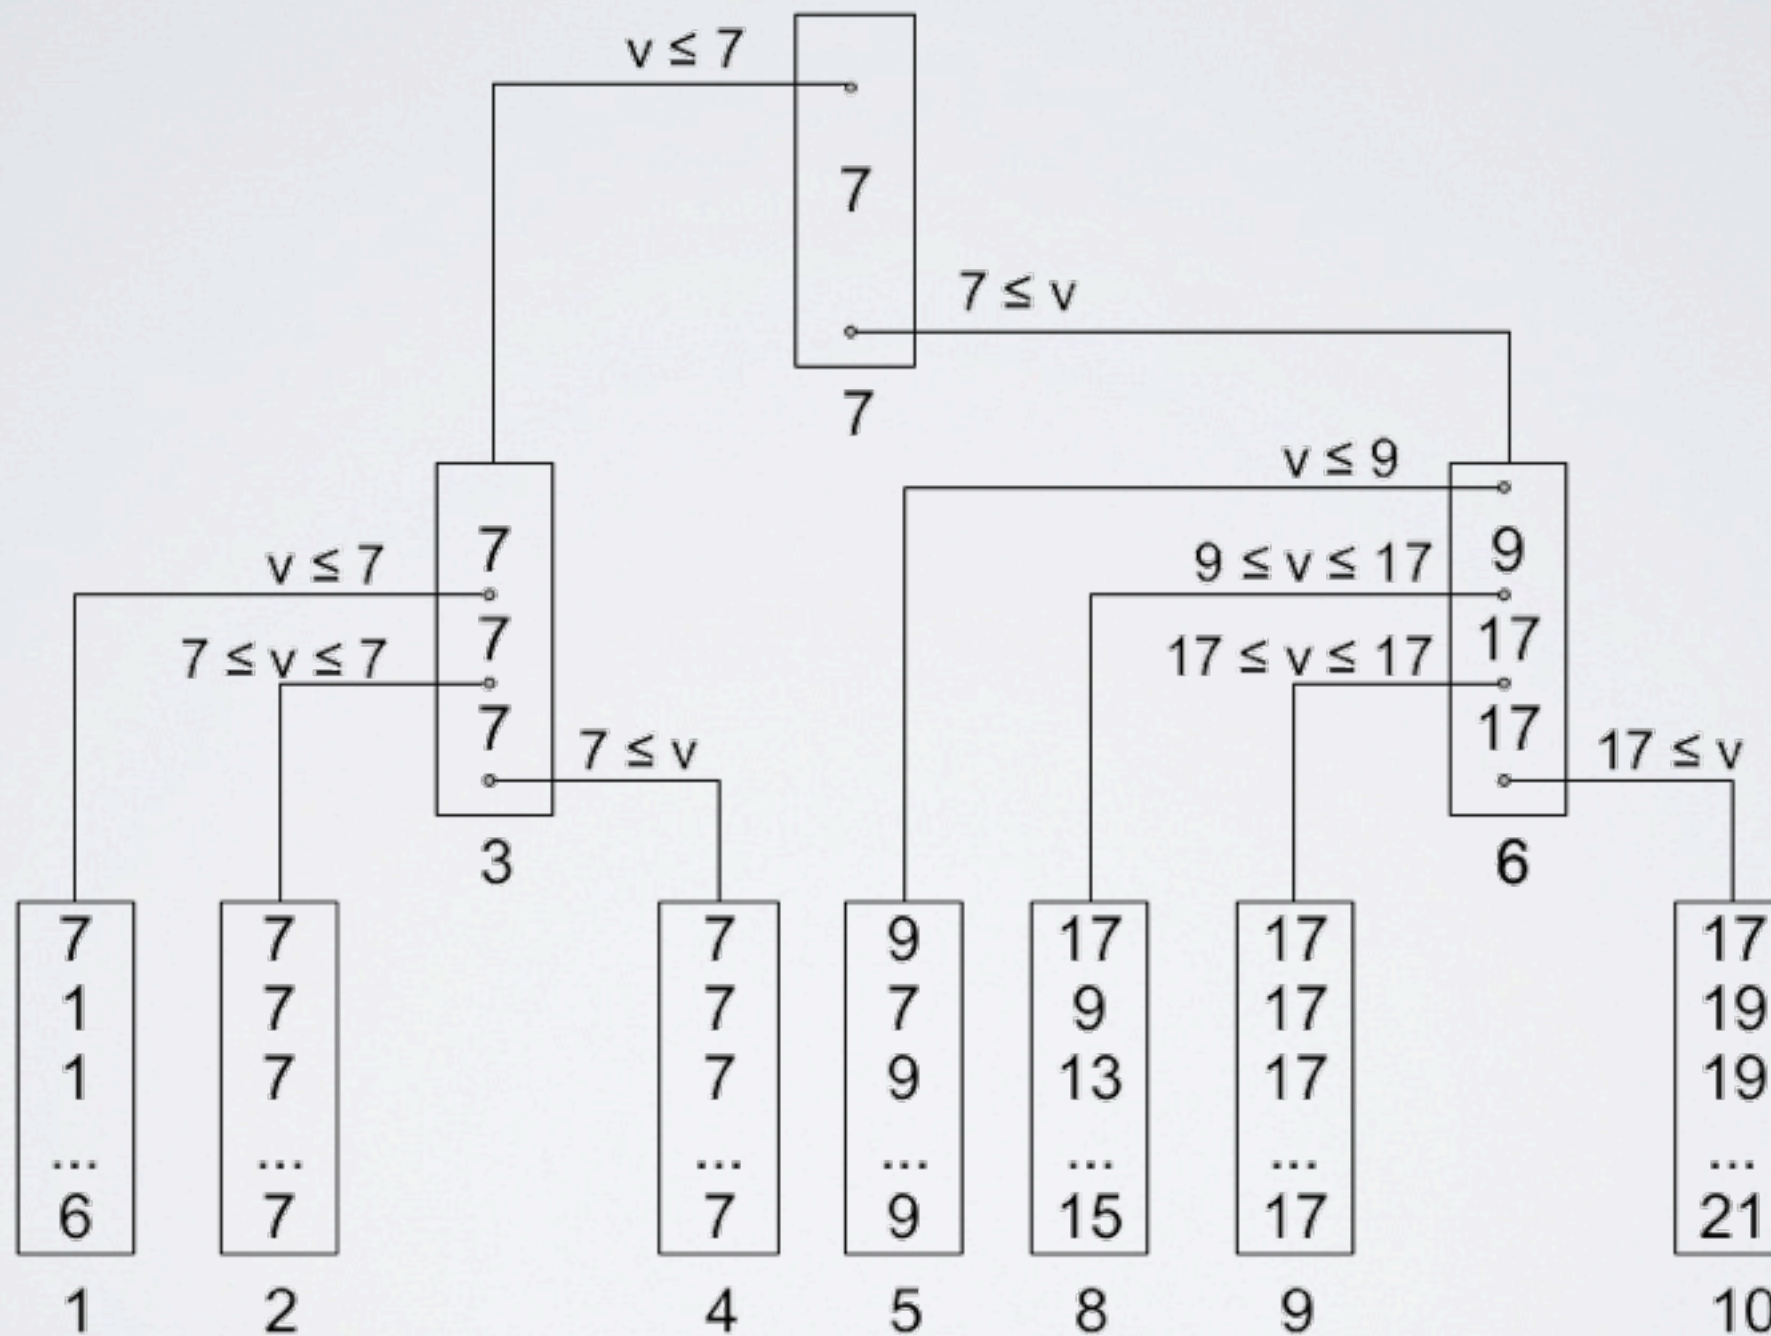
- Too much joins apply the brakes
- 3NF is a good start
- Wise Denormalisation



# B-TREE



# B-TREE





# FILLFACTOR

CREATE INDEX ... WITH (FILLFACTOR=n)

ALTER INDEX ... SET (FILLFACTOR=n)

$10 < n < 100$

# INDEX BLOAT

~~DROP INDEX name~~

~~CREATE INDEX name ON ...;~~

~~REINDEX name;~~

CREATE INDEX CONCURRENTLY name\_neu ON ...;

DROP INDEX name\_alt;


ALTER INDEX name\_neu RENAME TO name\_alt;



# INDEX USAGE

- Index on person(name, given\_name)  
... WHERE (name, given\_name)=('Miller', 'Solveig') ...
- Index on person(name) and Index on person(given\_name)  
... WHERE name='Miller' AND given\_name='Solveig' ...

# JOINS




SELECT \* FROM A JOIN B ON A.id=B.id;  
SELECT \* FROM A, B WHERE A.id=B.id;  
SELECT A.\* FROM A WHERE A.id IN  
(SELECT B.id FROM B);

INNER JOIN


## OUTER JOINS

---




LEFT JOIN

SELECT \* FROM A LEFT JOIN B  
ON A.id=B.id




RIGHT JOIN

SELECT \* FROM A RIGHT JOIN B  
ON B.id=A.id




FULL JOIN


SELECT \* FROM A FULL JOIN B  
ON A.id=B.id



WHERE B.id IS NULL



WHERE A.id IS NULL



WHERE A.id IS NULL OR B.id IS NULL



# CROSS JOINS

Each element with each element ....

Table A: (1,a), (3,b), (5,c)

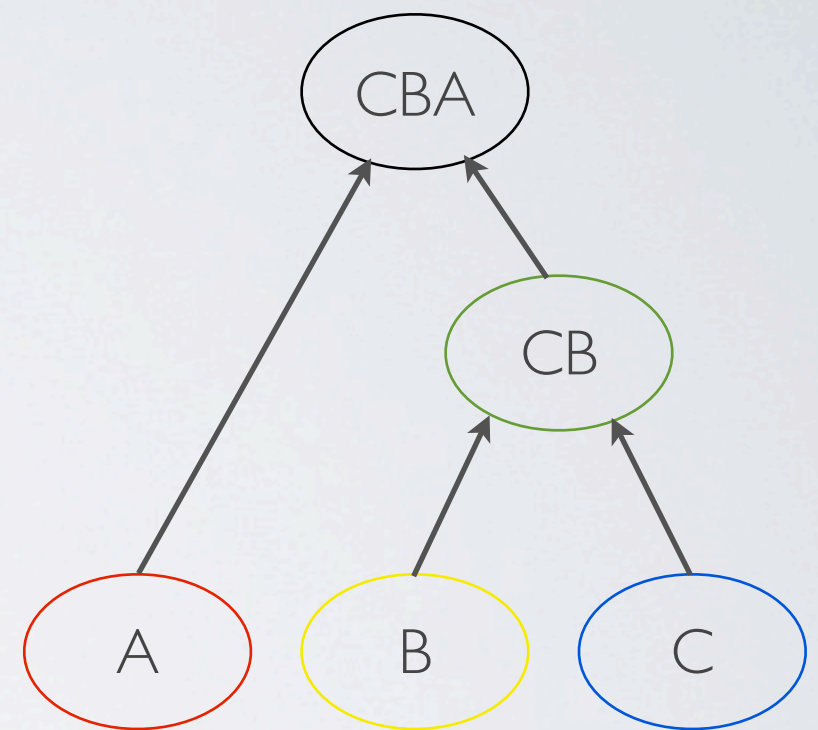


Table B: (1,x), (2,y), (3,z)

```
knolle=# SELECT * FROM a CROSS JOIN b;
```

i_a	wert_a	i_b	wert_b
1	a	1	x
1	a	2	y
1	a	3	z
3	b	1	x
3	b	2	y
3	b	3	z
5	c	1	x
5	c	2	y
5	c	3	z

(9 rows)

# WHAT'S CORRECT?





# WHAT'S FASTER?

```
SELECT * FROM a  
WHERE a.id NOT IN (SELECT id FROM b);
```

```
SELECT a.* FROM a LEFT JOIN b ON a.id=b.id  
WHERE b.id IS NULL;
```

# CORRELATED SUBSELECT?

```
SELECT town, violation, amount
FROM ticket_officer AS toff1
WHERE amount = (
    SELECT max(amount)
    FROM ticket_officer AS toff2
    WHERE toff2.town=toff1.town
);
```

=> needs 250 ms



# CTE

## Common Table Expression

```
WITH max_by_town AS (  
  SELECT town, max(amount) as total  
  FROM ticket_office  
  GROUP BY town  
)  
SELECT town, violation, amount  
FROM ticket_office  
WHERE town, amount IN (  
  SELECT town, total  
  FROM max_by_town  
);
```

=> needs 3 ms

# CTE

## Common Table Expression

```
WITH income AS (  
    SELECT town, sum(amount) AS total  
    FROM ticket_office  
    GROUP BY town  
)  
top AS (  
    SELECT town  
    FROM income  
    WHERE total > (SELECT avg(total) FROM income)  
)  
SELECT town, violation, sum(quantity) AS quantity, sum(amount) AS total  
FROM ticket_office  
WHERE town in (SELECT town FROM top)  
GROUP BY town, ticket;
```



# CTE

## Common Table Expression

```
WITH RECURSIVE meine(n) AS  
(  
    VALUES(1)  
    UNION ALL  
    SELECT n+1 FROM meine WHERE n < 100  
)  
SELECT SUM(n) FROM meine;
```

# PLANNER

- EXPLAIN - strategy
- EXPLAIN ANALYZE - strategy and execution



# EXPLAIN

```
ticket=# EXPLAIN SELECT t.town, tckt.violation, SUM(to.amount) AS total
FROM town AS t JOIN ticket_office AS to ON t.shortcut=to.town JOIN ticket AS tckt ON to.violation=tckt.violation
GROUP BY t.town, tckt.violation ORDER BY total DESC LIMIT 10;
```

## QUERY PLAN

---

```
Limit (cost=4878.07..4878.10 rows=10 width=69)
-> Sort (cost=4878.07..4941.18 rows=25245 width=69)
    Sort Key: (sum(to.amount))
-> GroupAggregate (cost=3827.64..4332.54 rows=25245 width=69)
    -> Sort (cost=3827.64..3890.75 rows=25245 width=69)
        Sort Key: t.town, tckt.violation
    -> Merge Join (cost=561.98..945.76 rows=25245 width=69)
        Merge Cond: (tckt.violation = to.violation)
    -> Sort (cost=71.17..73.72 rows=1020 width=32)
        Sort Key: tckt.violation
    -> Seq Scan on ticket tckt (cost=0.00..20.20 rows=1020 width=32)
-> Sort (cost=490.81..503.19 rows=4950 width=67)
    Sort Key: to.violation
-> Hash Join (cost=33.50..187.05 rows=4950 width=67)
    Hash Cond: ((t.shortcut)::text = (to.town)::text)
    -> Seq Scan on town t (cost=0.00..19.90 rows=990 width=48)
    -> Hash (cost=21.00..21.00 rows=1000 width=37)
        -> Seq Scan on ticket_office to (cost=0.00..21.00 rows=1000 width=37)
```

# EXPLAIN

- cost = estimated operation time; from\_ms .... to\_ms
- row = estimated number of found rows
- width = estimated row width given in Byte



# EXPLAIN ANALYZE

```
ticket=# EXPLAIN ANALYZE SELECT t.town, tckt.violation, SUM(to.amount) AS total
FROM town AS t JOIN ticket_office AS to ON t.shortcut=to.town JOIN ticket AS tckt ON to.violation=tckt.violation
GROUP BY t.town, tckt.violation ORDER BY total DESC LIMIT 10;
```

## QUERY PLAN

---

```
Limit (cost=4878.07..4878.10 rows=10 width=69) (actual time=26.814..26.815 rows=10 loops=1)
-> Sort (cost=4878.07..4941.18 rows=25245 width=69) (actual time=26.812..26.812 rows=10 loops=1)
    Sort Key: (sum(to.amount))
    Sort Method: top-N heapsort  Memory: 25kB
-> GroupAggregate (cost=3827.64..4332.54 rows=25245 width=69) (actual time=25.631..26.597 rows=256 loops=1)
    -> Sort (cost=3827.64..3890.75 rows=25245 width=69) (actual time=25.617..25.712 rows=1000 loops=1)
        Sort Key: t.town, tckt.violation
        Sort Method: quicksort  Memory: 125kB
    -> Merge Join (cost=561.98..945.76 rows=25245 width=69) (actual time=10.094..12.171 rows=1000 loops=1)
        Merge Cond: (tckt.violation = to.violation)
        -> Sort (cost=71.17..73.72 rows=1020 width=32) (actual time=0.102..0.103 rows=13 loops=1)
            Sort Key: tckt.violation
            Sort Method: quicksort  Memory: 25kB
        -> Seq Scan on ticket tckt (cost=0.00..20.20 rows=1020 width=32) (actual time=0.009..0.014 rows=13 loops=1)
    -> Sort (cost=490.81..503.19 rows=4950 width=67) (actual time=9.986..10.061 rows=1000 loops=1)
        Sort Key: to.violation
        Sort Method: quicksort  Memory: 125kB
    -> Hash Join (cost=33.50..187.05 rows=4950 width=67) (actual time=1.684..2.487 rows=1000 loops=1)
        Hash Cond: ((t.shortcut)::text = (to.town)::text)
        -> Seq Scan on town t (cost=0.00..19.90 rows=990 width=48) (actual time=0.003..0.011 rows=21 loops=1)
        -> Hash (cost=21.00..21.00 rows=1000 width=37) (actual time=1.659..1.659 rows=1000 loops=1)
            Buckets: 1024  Batches: 1  Memory Usage: 69kB
            -> Seq Scan on ticket_office to (cost=0.00..21.00 rows=1000 width=37) (actual time=0.007..0.674 rows=1000
loops=1)
Total runtime: 26.920 ms
```

# AFTER RUNNING ANALYZE

```
ticket=# EXPLAIN ANALYZE SELECT t.town, tckt.violation, SUM(to.amount) AS total
FROM town AS t JOIN ticket_office AS to ON t.shortcut=to.town JOIN ticket AS tckt ON to.violation=tckt.violation
GROUP BY t.town, tckt.violation ORDER BY total DESC LIMIT 10;
```

## QUERY PLAN

---

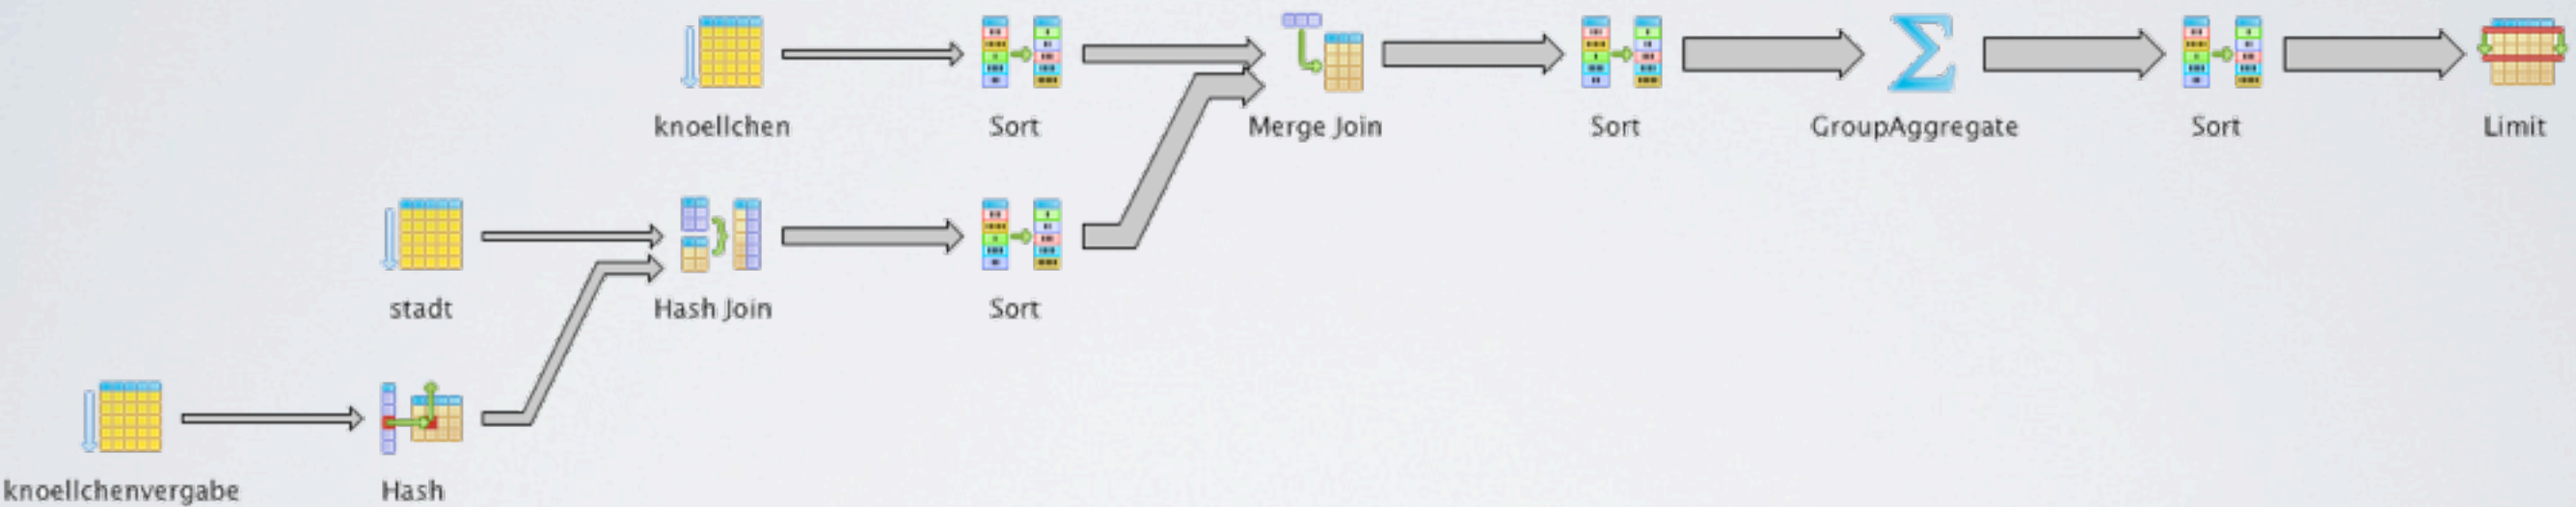
```
Limit (cost=67.39..67.42 rows=10 width=44) (actual time=5.586..5.590 rows=10 loops=1)
-> Sort (cost=67.39..68.08 rows=273 width=44) (actual time=5.584..5.586 rows=10 loops=1)
    Sort Key: (sum(to.amount))
    Sort Method: top-N heapsort  Memory: 25kB
-> HashAggregate (cost=58.77..61.49 rows=273 width=44) (actual time=5.080..5.240 rows=256 loops=1)
    -> Hash Join (cost=2.77..51.27 rows=1000 width=44) (actual time=0.084..2.812 rows=1000 loops=1)
        Hash Cond: (to.violation = tckt.violation)
        -> Hash Join (cost=1.47..36.22 rows=1000 width=44) (actual time=0.048..1.716 rows=1000 loops=1)
            Hash Cond: ((to.town)::text = (t.shortcut)::text)
            -> Seq Scan on ticket_office to (cost=0.00..21.00 rows=1000 width=37) (actual time=0.008..0.326 rows=1000 loops=1)
            -> Hash (cost=1.21..1.21 rows=21 width=12) (actual time=0.028..0.028 rows=21 loops=1)
                Buckets: 1024  Batches: 1  Memory Usage: 1kB
                -> Seq Scan on town t (cost=0.00..1.21 rows=21 width=12) (actual time=0.003..0.014 rows=21 loops=1)
        -> Hash (cost=1.13..1.13 rows=13 width=30) (actual time=0.027..0.027 rows=13 loops=1)
            Buckets: 1024  Batches: 1  Memory Usage: 1kB
            -> Seq Scan on ticket tckt (cost=0.00..1.13 rows=13 width=30) (actual time=0.008..0.015 rows=13 loops=1)
Total runtime: 5.686 ms
```



# ACTUAL

- time = needed operation time; from\_ms .... to\_ms
- row = Number of found rows
- loops = Number of executions per operation

# PGADMIN III





# BREAKDOWN

- (cost=0.00..19.90 rows=990 width=48) (actual time=0.003..0.011 rows=21 loops=1)
  - ANALYZE or STATISTIC TARGET
- (actual time=10.081..15.764 rows=1000 loops=651)
  - Think about logic, Redesign Query, CTE (Common Table Expression)
- (actual time=25.617..12425.712 rows=1000 loops=1)
  - Think about logic, Redesign Query, Indexes

# STATISTIC

- Random Sample
- postgresql.conf: default\_statistic\_target = 100
- ALTER TABLE ... ALTER COLUMN ... SET STATISTIC value;



# ANALYZE

- SQL Command
- PostgreSQL also allows British: ANALYSE
- Frequency analysis
- Statistic tables like pg\_class
- autovacuum includes autoanalyze

# EXPLAIN.DEPEZ.COM

options						
HTML	TEXT	STATS				
#	exclusive	inclusive	rows x	rows	loops	node
1.	0.003	26.815	↑ 1.0	10	1	→ Limit (cost=4878.07..4878.10 rows=10 width=69) (actual time=26.814..26.815 rows=10 loops=1)
2.	0.215	26.812	↑ 2524.5	10	1	→ Sort (cost=4878.07..4941.18 rows=25245 width=69) (actual time=26.812..26.812 rows=10 loops=1) Sort Key: (sum(kv.betrag)) Sort Method: top-N heapsort Memory: 25kB
3.	0.885	26.597	↑ 98.6	256	1	→ GroupAggregate (cost=3827.64..4332.54 rows=25245 width=69) (actual time=25.631..26.597 rows=256 loops=1)
4.	13.541	25.712	↑ 25.2	1000	1	→ Sort (cost=3827.64..3890.75 rows=25245 width=69) (actual time=25.617..25.712 rows=1000 loops=1) Sort Key: s.stadt, k.verstoss Sort Method: quicksort Memory: 125kB
5.	2.007	12.171	↑ 25.2	1000	1	→ Merge Join (cost=561.98..945.76 rows=25245 width=69) (actual time=10.094..12.171 rows=1000 loops=1) Merge Cond: (k.verstoss = kv.verstoss)
6.	0.089	0.103	↑ 78.5	13	1	→ Sort (cost=71.17..73.72 rows=1020 width=32) (actual time=0.102..0.103 rows=13 loops=1) Sort Key: k.verstoss Sort Method: quicksort Memory: 25kB
7.	0.014	0.014	↑ 78.5	13	1	→ Seq Scan on knoelchen k (cost=0.00..20.20 rows=1020 width=32) (actual time=0.009..0.014 rows=13 loops=1)
8.	7.574	10.061	↑ 5.0	1000	1	→ Sort (cost=490.81..503.19 rows=4950 width=67) (actual time=9.986..10.061 rows=1000 loops=1) Sort Key: kv.verstoss Sort Method: quicksort Memory: 125kB
9.	0.817	2.487	↑ 5.0	1000	1	→ Hash Join (cost=33.50..187.05 rows=4950 width=67) (actual time=1.684..2.487 rows=1000 loops=1) Hash Cond: ((s.kennzeichen)::text = (kv.stadt)::text)
10.	0.011	0.011	↑ 47.1	21	1	→ Seq Scan on stadt s (cost=0.00..19.90 rows=990 width=48) (actual time=0.003..0.011 rows=21 loops=1)
11.	0.985	1.659	↑ 1.0	1000	1	→ Hash (cost=21.00..21.00 rows=1000 width=37) (actual time=1.659..1.659 rows=1000 loops=1) Buckets: 1024 Batches: 1 Memory Usage: 69kB
12.	0.674	0.674	↑ 1.0	1000	1	→ Seq Scan on knoelchenvergabe kv (cost=0.00..21.00 rows=1000 width=37) (actual time=0.007..0.674 rows=1000 loops=1)

Hubert Lubaczewski, Nickname: depesz



# EXPLAIN.DEPEZ.COM

HTML TEXT **STATS**

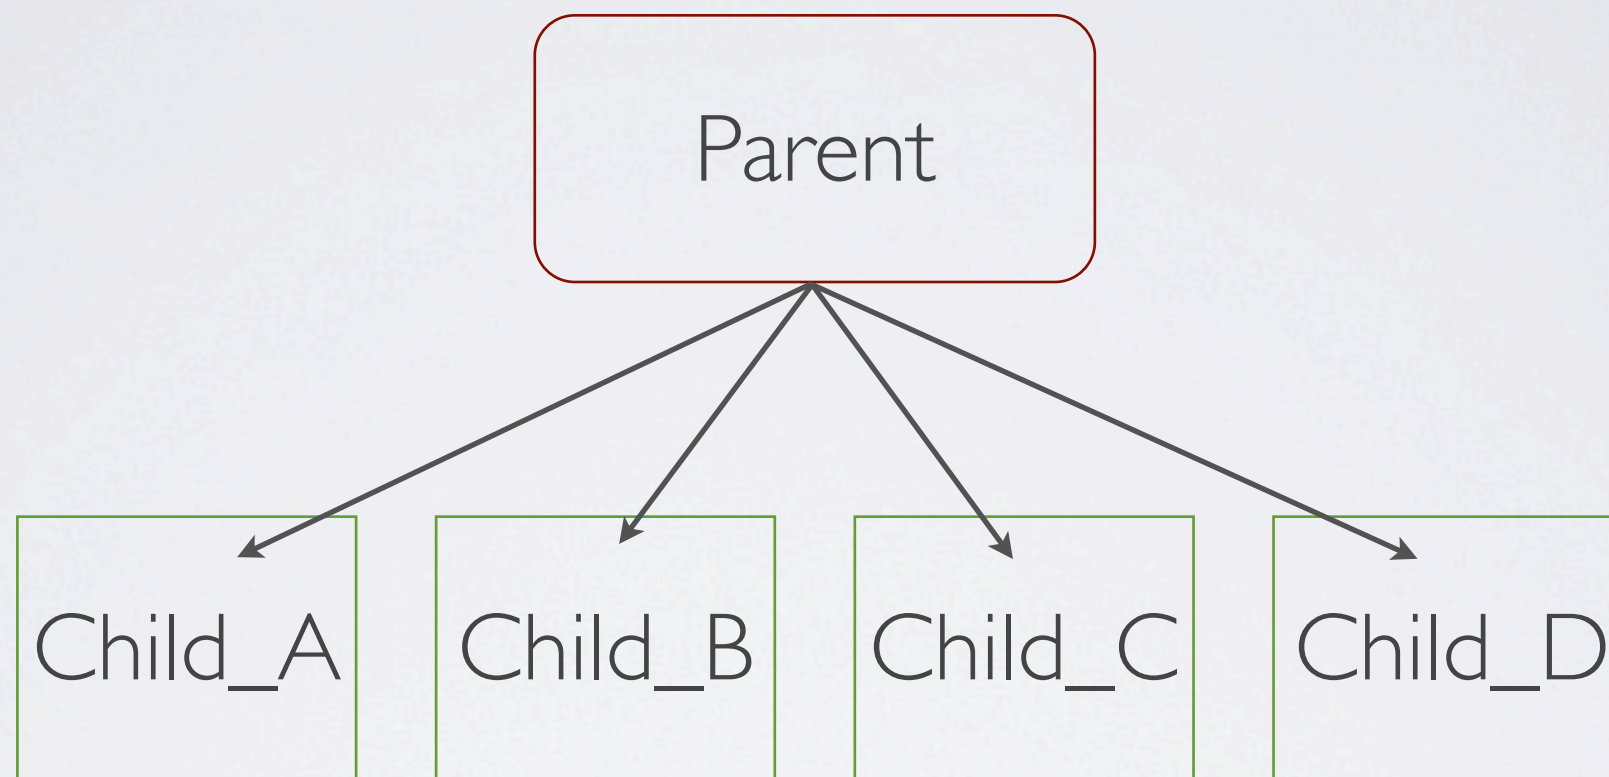
## Per node type stats

node type	count	sum of times	% of query
GroupAggregate	1	0.885 ms	3.3 %
Hash	1	0.985 ms	3.7 %
Hash Join	1	0.817 ms	3.0 %
Limit	1	0.003 ms	0.0 %
Merge Join	1	2.007 ms	7.5 %
Seq Scan	3	0.699 ms	2.6 %
Sort	4	21.419 ms	79.9 %

## Per table stats

Table name	Scan count	Total time	% of query
scan type	count	sum of times	% of table
<b>knoellchen</b>	1	0.014 ms	0.1 %
Seq Scan	1	0.014 ms	100.0 %
<b>knoellchenvergabe</b>	1	0.674 ms	2.5 %
Seq Scan	1	0.674 ms	100.0 %
<b>stadt</b>	1	0.011 ms	0.0 %
Seq Scan	1	0.011 ms	100.0 %

# PARTITIONING





# THANKS FOR LISTENING